

HPC HW 0 Problem 2

Noah D. Brenowitz

February 9, 2015

The source code for the Gauss-Seidel method with over-relaxation is give in the file `laplace.c`:

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>

#define STOP_ITER_RAT 10e-6
#define OMEGA_OVER_RELAXED 1.5

double calc_resid(int N, double* f, double *u){
    int i, j;
    double resid = 0.0;
    double ai;

    double h2 = 1.0/(N+1)/(N+1);

    ai = (2 * u[0] - u[1])/h2 -f[0];
    resid += ai*ai;

    ai = (2 * u[N-1] - u[N-2])/h2 -f[N-1];
    resid += ai*ai;

    for (i = 1; i < N-1; i++) {
        ai = (- u[i+1] + 2 * u[i] - u[i-1])/h2 -f[i];
        resid += ai*ai;
    }

    return sqrt(resid);
}

void gauss_seidel_laplace(int N, double *f, double *u){
    int i;
    double h2 = 1.0/(N+1)/(N+1);

    u[0] = (h2 * f[0] + u[1])/2.0;

    for (i = 1; i < N-1; i++) {
```

```

    u[i] = (h2 * f[i] + u[i-1] + u[i+1])/2.0;
}

u[N-1] = (h2* f[N-1] + u[N-2])/2.0;
}

void overrelaxed_gauss_seidel_laplace(int N, double *f, double *u, double* w){
    int i;
    double h2 = 1.0/(N+1)/(N+1);

    for (i = 0; i < N; i++)
        w[i] = u[i];

    u[0] = (h2 * f[0] + u[1])/2.0;

    for (i = 1; i < N-1; i++) {
        u[i] = (h2 * f[i] + u[i-1] + u[i+1])/2.0;
    }

    u[N-1] = (h2* f[N-1] + u[N-2])/2.0;

    for (i = 0; i < N; i++)
        u[i] = (1-OMEGA_OVER_RELAXED) * w[i] + OMEGA_OVER_RELAXED * u[i];
}

void test_resid(){
    int N = 100;

    double *u, *f;

    // allocate arrays
    u = (double *) malloc(N*sizeof (double));
    f = (double *) malloc(N*sizeof (double));

    // initialize f and u
    int i;
    for (i = 0; i < N; i++) {
        f[i] = 2.0;
        u[i] = 0.0;
    }

    double resid;

    resid = calc_resid(N, f, u);
    printf("Resid is %f", resid);
}

```

```

}

int main(int argc, char *argv[])
{
    const int run_over_relaxed = 0; // if not 0 run with overrelaxed version of algorithm

    int N = atoi(argv[1]);

    double *u, *f, *w;

    // allocate arrays
    u = (double *) malloc(N*sizeof (double));
    f = (double *) malloc(N*sizeof (double));

    if (run_over_relaxed){
        w = (double *) malloc(N*sizeof (double));
    }

    // initialize f and u
    int i;
    for (i = 0; i < N; i++) {
        f[i] = 1.0;
        u[i] = 0.0;
    }

    // Begin iterations
    double resid_init, resid_cur;
    resid_init = calc_resid(N, f, u);
    resid_cur = resid_init;

    while (resid_cur / resid_init > STOP_ITER_RAT){

        if (run_over_relaxed){
            overrelaxed_gauss_seidel_laplace(N, f, u, w);
        } else {
            gauss_seidel_laplace(N, f, u);
        }

        resid_cur = calc_resid(N, f, u);
        printf("Resid is %f\n", resid_cur );
    }

    // deallocate
    free(f);
    free(u);
}

```

```

    return 0;
}

```

This source is compiled and timed using the following shell script (`run.sh`):

```

CC=gcc
CFLAGS=-lm

${CC} ${CFLAGS} -o laplace laplace.c
${CC} ${CFLAGS} -O0 -o laplace0 laplace.c
${CC} ${CFLAGS} -O3 -o laplace3 laplace.c

```

```

N=1000
time ./laplace $N > /dev/null
time ./laplace0 $N >/dev/null
time ./laplace3 $N > /dev/null

```

The output for $N=1000$ is:

```

$ zsh run.sh
./laplace $N > /dev/null 12.03s user 0.00s system 100% cpu 12.031 total
./laplace0 $N > /dev/null 12.05s user 0.00s system 100% cpu 12.045 total
./laplace3 $N > /dev/null 9.48s user 0.00s system 100% cpu 9.477 total

```

There is around a 30% speed-up for the `-O3` optimized file compared to the non-optimized version. The number of iterations required to reach convergence is

```

./laplace 1000 | wc -l

1158227

```

For $N=100$, the output is:

```

zsh run.sh
./laplace $N > /dev/null 0.02s user 0.00s system 94% cpu 0.017 total
./laplace0 $N > /dev/null 0.01s user 0.00s system 98% cpu 0.016 total
./laplace3 $N > /dev/null 0.01s user 0.00s system 92% cpu 0.013 total

```

So the $N=100$ code runs much more quickly. The number of iterations required for convergence here is

```

./laplace 100 | wc -l

11796

```